



Building a package and submit as an application

Right click on the project → Run As → Maven Clean

Right click on the project → Run As → Maven Install

Goto Linux terminal

`cd ~`

Submit Jar to different clusters:-

Local Mode

```
spark-submit --class org.inceptez.streaming.filestream --master local[2]  
/home/hduser/workspacespark/spark/target/spark-0.0.1-SNAPSHOT-jar-with-dependencies.jar
```

Standalone Mode

```
start-master.sh  
start-slaves.sh
```

Submit the jar in standalone mode

```
spark-submit --class org.inceptez.streaming.filestream --master spark://localhost:7077  
/home/hduser/workspacespark/spark/target/spark-0.0.1-SNAPSHOT-jar-with-dependencies.jar
```

Yarn mode

```
spark-submit --class org.inceptez.streaming.kafkastream --master yarn  
/home/hduser/workspacespark/spark/target/spark-0.0.1-SNAPSHOT-jar-with-dependencies.jar
```

```
spark-submit --class org.inceptez.streaming.filestream --master yarn  
/home/hduser/workspacespark/spark/target/spark-0.0.1-SNAPSHOT.jar --driver-memory  
512m --num-executors 1 --executor-cores 2 --executor-memory 512m
```

```
spark-submit --class org.inceptez.stream.filestream --master yarn  
/home/hduser/install/Sparktwitterusecase/target/sparkIncetez-0.0.1-SNAPSHOT.jar --driver-  
memory 512m --num-executors 1 --executor-cores 1 --executor-memory 1 --spark-shuffle-  
compress true --spark-speculation true --spark-dynamicAllocation-enabled true --spark-  
dynamicAllocation-minExecutors 1 --spark-dynamicAllocation-maxExecutors 1 --spark-  
dynamicAllocation-initialExecutors 1
```

```
spark-submit --verbose \  
--class org.inceptez.stream.filestream \  
--master yarn /home/hduser/install/Sparktwitterusecase/target/sparkIncetez-0.0.1-  
SNAPSHOT.jar \  
--deploy-mode client \  
--queue default \  
--driver-memory 512m \  
--num-executors 1 \  
--executor-cores 2 \  
--executor-memory 512m \  
--spark-shuffle-compress true \  
--spark-speculation true \  
--spark-dynamicAllocation-enabled true \  
--spark-dynamicAllocation-initialExecutors 1 \  
--spark-dynamicAllocation-minExecutors 1 \  
--spark-dynamicAllocation-maxExecutors 8 \  
--spark-dynamicAllocation-executorIdleTimeout 30s \  
--conf spark.shuffle.service.enabled true \  
--conf spark.sql.shuffle.partitions=4
```

<https://spark.apache.org/docs/2.0.1/configuration.html>

To Start spark shell with yarn-client

```
spark-shell --verbose --master yarn-client --driver-memory 512m --num-executors 1 --  
executorcores 2 --executor-memory 512m
```

or

```
spark-shell --master yarn
```

Spark Memory Allocation

Let's consider a 10 node cluster with following config and analyse different possibilities of executors-core-memory distribution:

****Cluster Config:****

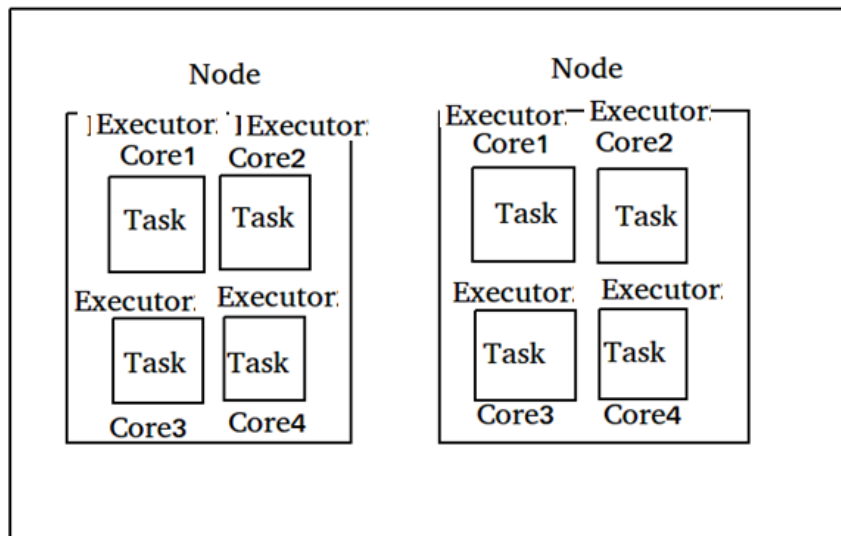
12 node (10 Worker/Data Nodes, 1 Master, 1 Client)

16 cores per Node

64GB RAM per Node

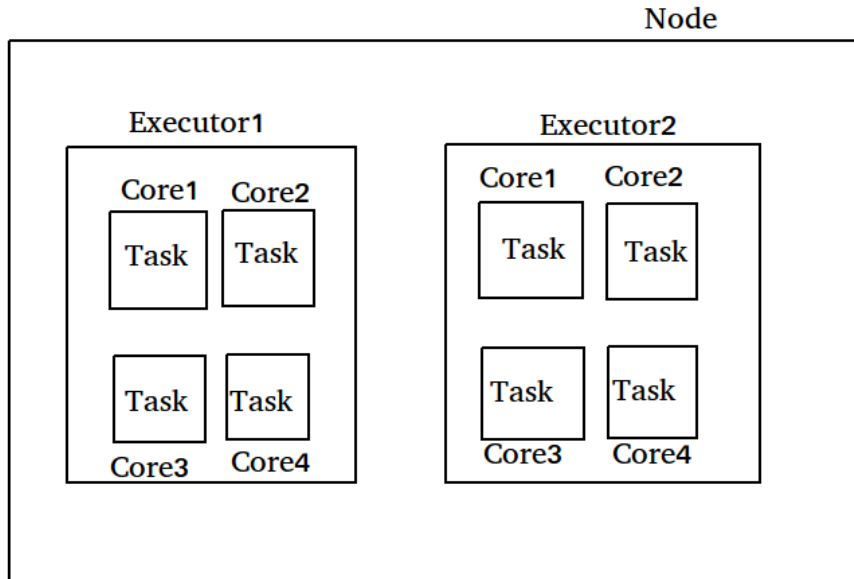
First Approach: Tiny executors [One Executor per core]:

1. Multiple tasks in the same executor JVM can't be run, huge number of small JVMs start and kill overhead.
2. Shared/cached variables like broadcast variables and accumulators will be replicated in each executor of the nodes which is 16 times.
3. Not leaving enough memory overhead for Hadoop/Yarn daemon processes and not counting in ApplicationManager.



Second Approach: Fat executors (One Executor per node)

1. With all 16 cores per executor,
2. ApplicationMaster and daemon processes are not counted
3. HDFS throughput will hurt and it'll result in excessive garbage results.



Third Approach: Balance between Fat (vs) Tiny

With the cluster of 10 Nodes , 16 cores per Node, 64GB RAM per Node

Assign 5 cores per executors => `--executor-cores = 5` (for good HDFS throughput)

Leave 1 core per node for DN/Yarn daemons => Num cores available per node = $16 - 1 = 15$

Total available of cores in cluster = $15 \text{ cores} \times 10 \text{ nodes} = 150 \text{ cores}$

Number of available executors = $(\text{total cores} / \text{num-cores-per-executor}) = 150 \text{ total cores} / 5 \text{ executor cores} = 30 \text{ executors available}$

Leave 1 executor for ApplicationMaster => $30 - 1 = 29$ total executors

Number of executors per node => `--num-executors = 30` executors / 10 nodes = `3 executors per node`

Memory per executor => $64\text{GB per node} / 3 \text{ executors} = 21\text{GB per executor}$

Memory per executor after Counting off heap overhead = off heap allocation of ~7% of 21GB = 3GB, hence `--executor-memory = 21 - 3 = 18GB` Memory per executor

So, recommended config is: 29 executors, 18GB memory each and 5 cores each!!

Third approach has found right balance between Fat vs Tiny approaches.

Finally it achieved parallelism of a fat executor and best throughputs of a tiny executor!!

Leaving all the above calculations, iterate or benchmark for optimal/better performances considering ondemand/continuous other resources consumption in the cluster for other jobs at that time.